

Istat

Deep Learning & NLP

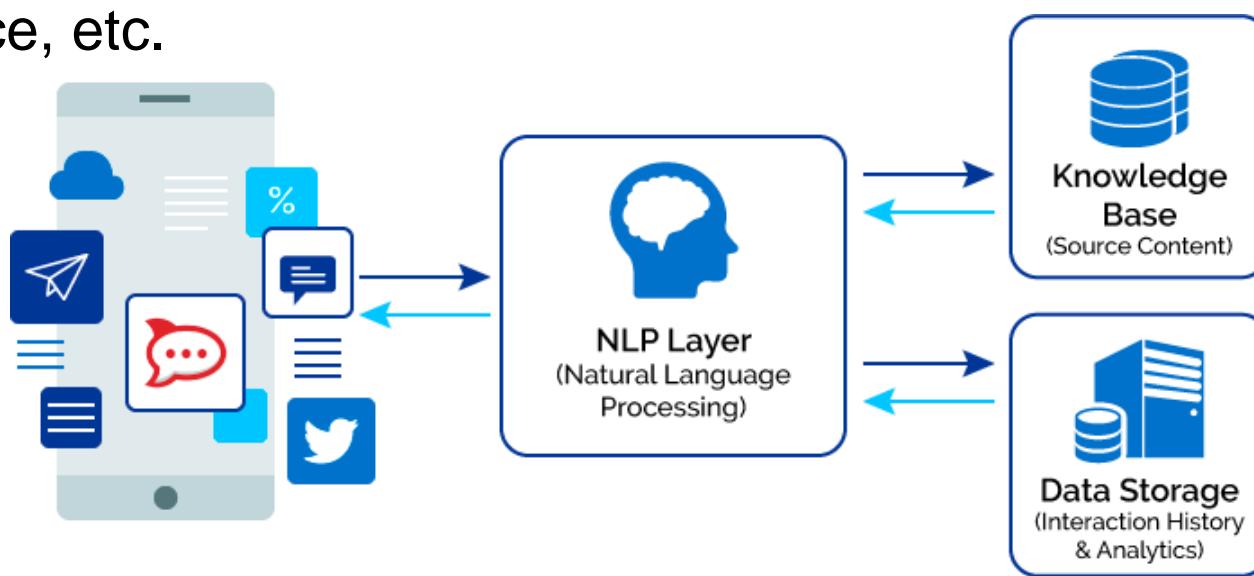
Francesco Pugliese, PhD

*Italian National Institute of Statistics, Division
"Information and Application Architecture", Directorate
for methodology and statistical design*

Email Francesco Pugliese : francesco.pugliese@istat.it

Textual Big Data alias The problem of the Natural Languale Processing - NLP

- Understanding **complex language utterances** is one of the **hardest challenge** for Artificial Intelligence (AI) and Machine Learning (ML).
- **NLP** is everywhere because people communicate most everything: web search, advertisement, emails, customer service, etc.



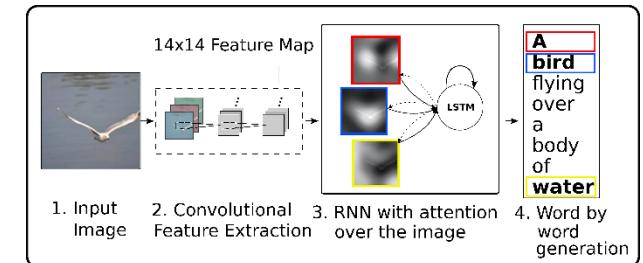
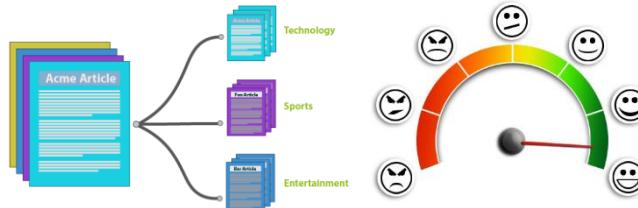
Deep Learning and NLP

- “Deep Learning” approaches have obtained very high performance across many different **NLP** tasks. These models can often be trained with a **single end-to-end model** and do not require traditional, task-specific feature engineering.
(Stanford University School Of Engineering – CS224D)
- **Natural language processing** is shifting from statistical methods to **Neural Networks**.



7 NLP applications where Deep Learning achieved «state-of-art» performance

- **1 Text Classification:** Classifying the topic or theme of a document (i.e. Sentiment Analysis).
- **2 Language Modeling:** Predict the **next word given the previous words**. It is fundamental for other tasks.
- **3 Speech Recognition:** Mapping an **acoustic signal** containing a spoken natural language utterance into the corresponding sequence of words intended by the speaker.
- **4 Caption Generation:** Given a **digital image**, such as a photo, generate a **textual description** of the contents of the image.



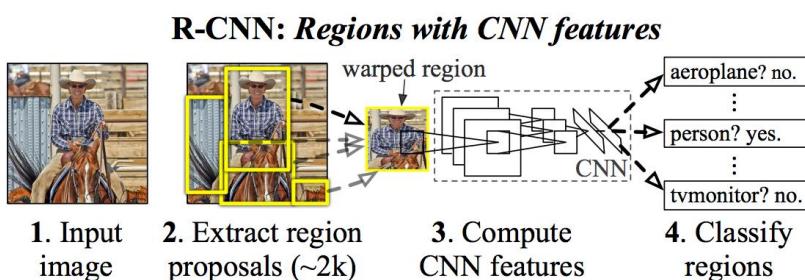
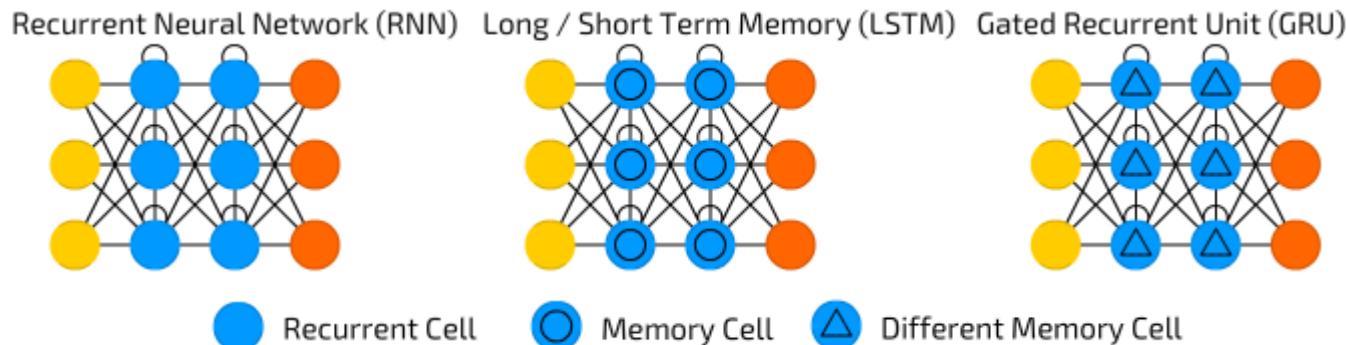
7 NLP applications where Deep Learning achieved «state-of-art» performance

- **5 Machine Translation:** Automatic translation of text or speech from one language to another, is one [of] the most important applications of NLP.
- **6 Document Summarization:** It is the task where a short description of a text document is created.
- **7 Question Answering:** It is the task where the system tries to answer a user query that is formulated in the form of a question by returning the appropriate noun phrase such as a location, a person, or a date. (i.e. Who killed President Kennedy? Oswald)



Text Classification Models

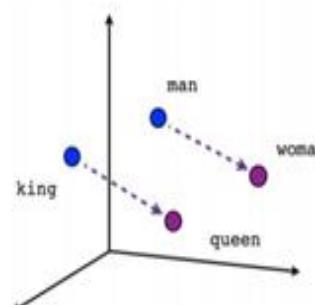
- **RNN, LSTM, GRU, ConvLstm, RecursiveNN, RNTN, RCNN**
- The modus operandi for text classification involves the use of a pre-trained **word embedding** for **representing words** and a **deep neural networks** for **learning how to discriminate documents** on classification problems.



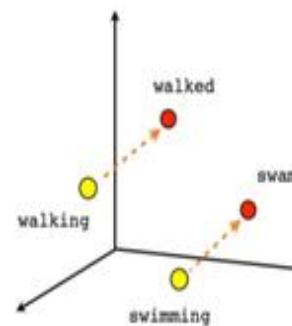
- The **non-linearity of the NN** leads to superior classification accuracy.

Word Embedding & Language Modeling

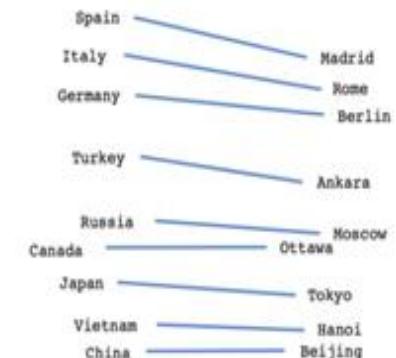
- Word embedding is the collective name for a set of language modeling and feature learning techniques for natural language processing (NLP) where words or sentences from the vocabulary are mapped to vectors of real numbers.
- These vectors are semantically correlated by metrics like cosine distance



Male-Female

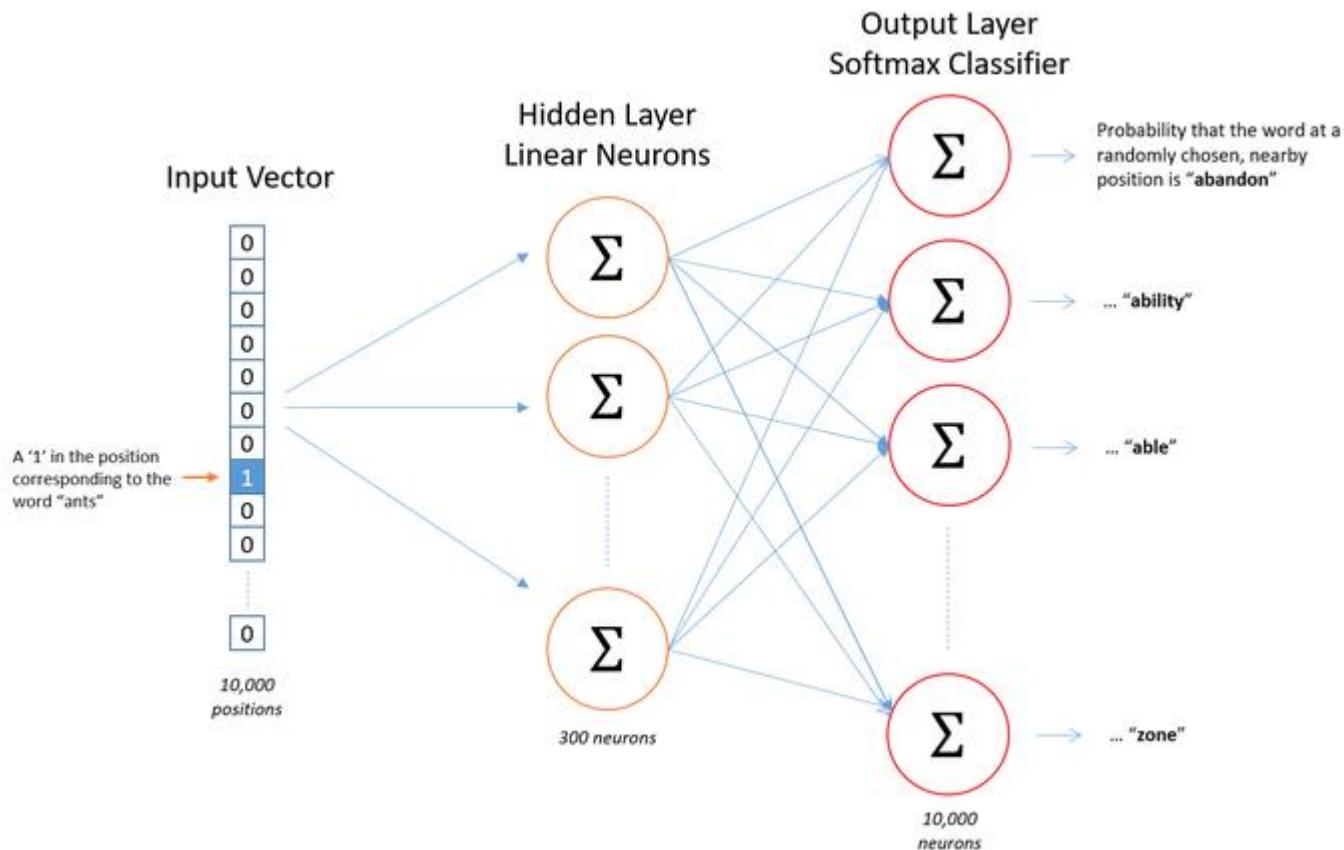


Verb tense

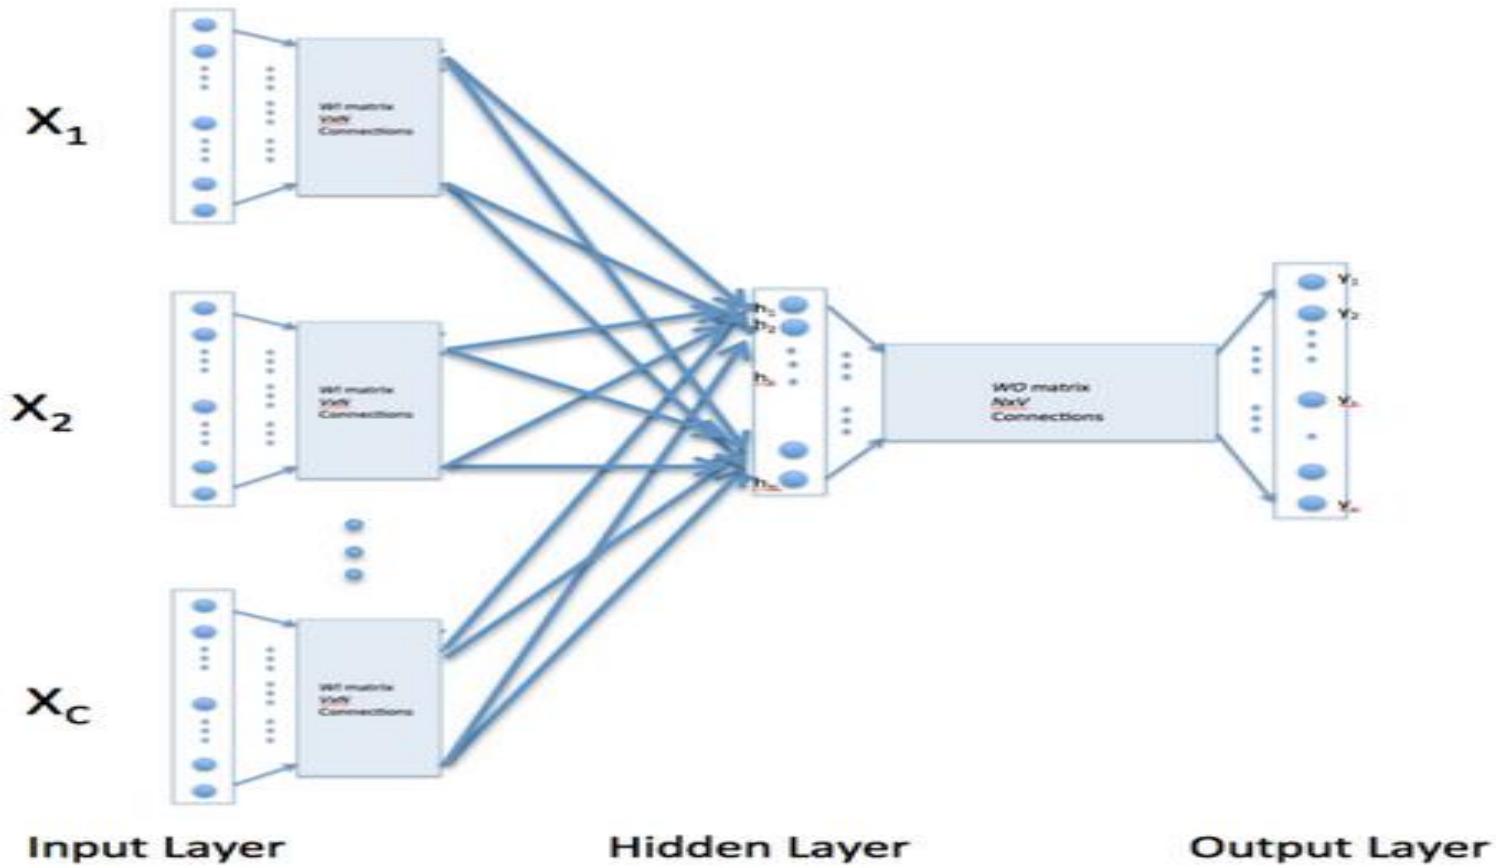


Country-Capital

Skip-Gram Model (Mikolov, et. al., 2013)



C-BOW Model (Bow, et al., 2003).



Sentiment Analysis (*Ain, et al. 2017*)

- **Sentiments** of users that are expressed on the web has great influence on the readers, product vendors and politicians.
- **Sentiment Analysis** refers to text organization for the classification of mind-set or feelings in different manners such as negative, positive, favorable, unfavorable, thumbs up, thumbs down, etc. Thanks to DL, the SA can be visual as well.



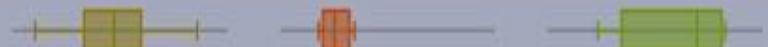
Discovering people opinions, emotions and feelings about
a product or service

Sentiment Analysis with Feedback

Stockle [start page](#)



Apple Inc. **AAPL** 116.30 (+0.25%)



ADBE **ADBE** 0.0 (0.0%)



eBay Inc. **EBAY** 31.46 (-0.49%)



GOOGL **GOOGL** 0.0 (0.0%)



Microsoft Corporation **MSFT** 57.19 (-0.85%)

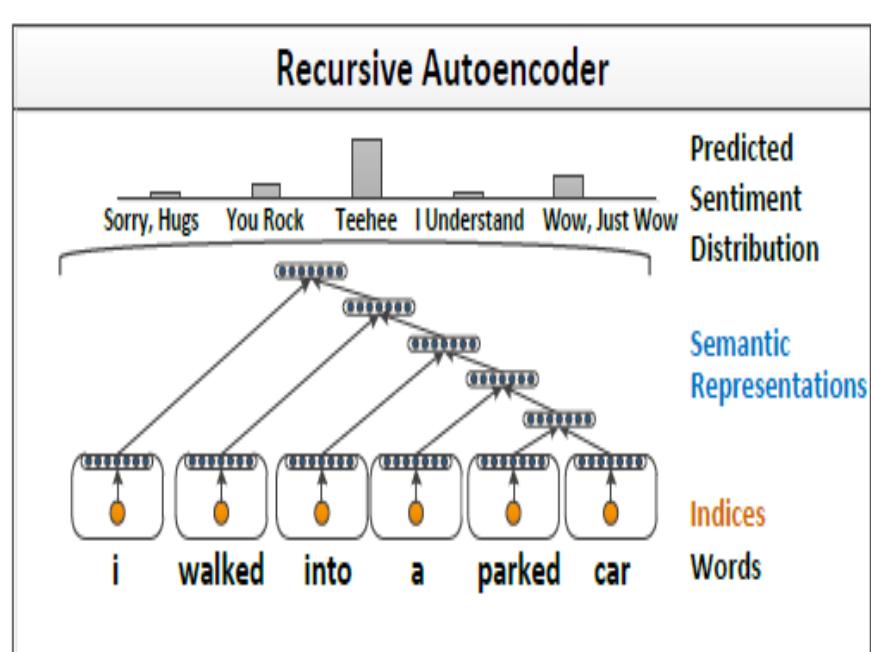


Yahoo! Inc. **YHOO** 42.68 (-1.24%)



Recursive Neural Tensor Networks (RecursiveNN) (Socher, R., et al., 2011b)

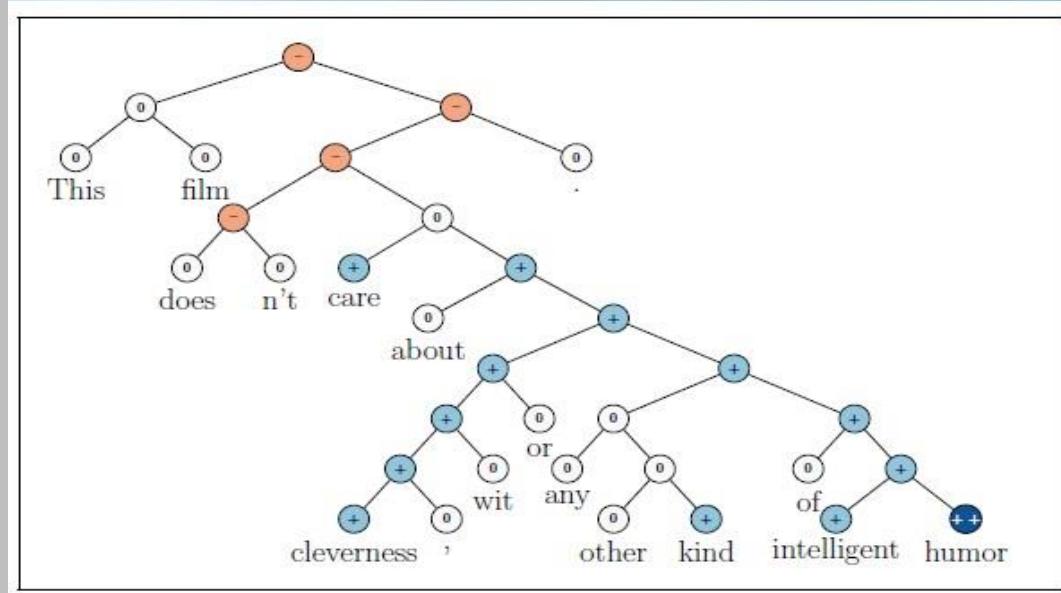
- This models are recursive auto-encoders which learn semantic vector representations of phrases. Word indices (orange) are first mapped into a semantic vector space (blue).
- Then they are recursively merged by the same auto-encoder network into a fixed length sentence representation. The vectors at each node are used as features to predict a distribution over text labels.



Recursive Neural Tensor Networks (RNTN)

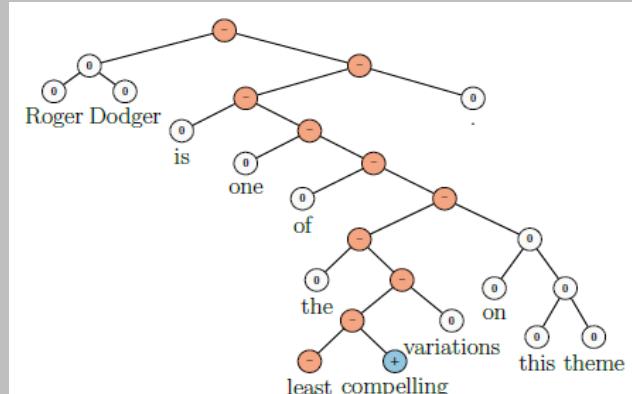
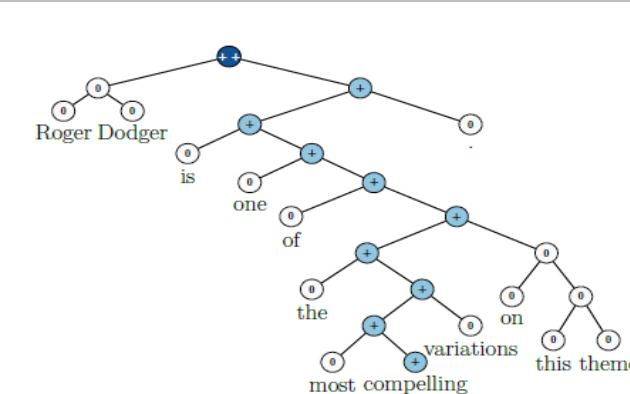
(Socher, R., et al. 2013)

- The Stanford Sentiment Treebank is the first corpus with fully labeled parse trees that allows for a complete analysis of the compositional effects of sentiment in language.
- RNTNs compute parent vectors in a bottom up fashion using a compositionality function and use node vectors as features for a classifier at that node.



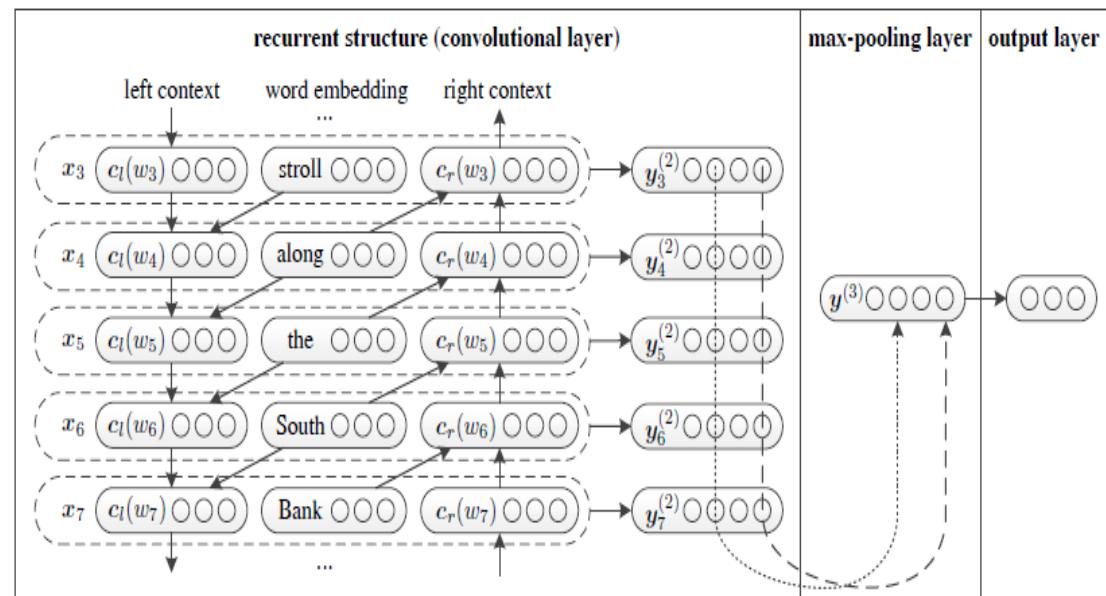
RNTN – Upside and Downside

- RNTNs are very efficient in terms of constructing sentence representations.
- RNTNs capture the semantics of a sentence via a tree structure. Its performance heavily depends on the performance of the textual tree construction.
- Constructing such a textual tree exhibits a time complexity of at least $O(n^2)$, where n is the length of the text.
- RNTNs are unsuitable for modeling long sentences or documents.



Recurrent Convolutional Neural Networks (RCNN) (Lai, S., et al. 2015)

- They adopt a recurrent structure to **capture contextual information** as far as possible when learning word representations, which may introduce considerably **less noise compared** to traditional window-based neural networks.
- The **bi-directional recurrent structure** of RCNNs.
- **RCNNs** exhibit a time complexity of $O(n)$



RCNN Equations

- RCNNs exhibit a **time complexity of $O(n)$** , which is linearly correlated with the length of the text length.

$$c_l(w_i) = f(W^{(l)} c_l(w_{i-1}) + W^{(sl)} e(w_{i-1})) \quad (1)$$

$$c_r(w_i) = f(W^{(r)} c_r(w_{i+1}) + W^{(sr)} e(w_{i+1})) \quad (2)$$

- **7 equations** defining all the Neural Network topology

$$x_i = [c_l(w_i); e(w_i); c_r(w_i)] \quad (3)$$

$$y_i^{(2)} = \tanh (W^{(2)} x_i + b^{(2)}) \quad (4)$$

$$y^{(3)} = \max_{i=1}^n y_i^{(2)} \quad (5)$$

- **Input length** can be variable

$$y^{(4)} = W^{(4)} y^{(3)} + b^{(4)} \quad (6)$$

$$p_i = \frac{\exp (y_i^{(4)})}{\sum_{k=1}^n \exp (y_k^{(4)})} \quad (7)$$

RCNN in Keras

```
class SentimentModelRecConvNet:  
    @staticmethod  
  
    def build(input_length, vector_dim):  
        hidden_dim_RNN = 200  
        hidden_dim_Dense = 100  
  
        embedding = Input(shape=(input_length, vector_dim))  
  
        left_context = LSTM(hidden_dim_RNN, return_sequences = True)(embedding) # Equation 1  
        # left_context: batch_size x tweet_length x hidden_state_dim  
        right_context = LSTM(hidden_dim_RNN, return_sequences = True, go_backwards = True)(embedding) # Equation 2  
        # right_context: come left_context  
        together = concatenate([left_context, embedding, right_context], axis = 2) # Equation 3  
        semantic = TimeDistributed(Dense(hidden_dim_Dense, activation = "tanh"))(together) # Equation 4  
        pool_rnn = Lambda(lambda x: backend.max(x, axis = 1), output_shape = (hidden_dim_Dense, ))(semantic) # Equation 5  
        pool_rnn_args = Lambda(lambda x: backend.argmax(x, axis=1), output_shape = (hidden_dim_Dense, ))(semantic)  
  
        output = Dense(1, input_dim = hidden_dim_Dense, activation = "sigmoid")(pool_rnn) # Equations 6, 7  
  
        deepnetwork = Model(inputs=embedding, outputs=output)  
        deepnetwork_keywords = Model(inputs=embedding, outputs=pool_rnn_args)  
  
        return [deepnetwork, deepnetwork.keywords]
```

RCNN: Feature Extraction

- RCNNs employ a max-pooling layer that automatically judges which words play key roles in text classification to capture the key components in texts.
- The most important words are the information most frequently selected in the max-pooling layer.
- Contrary to the most positive and most negative phrases in RNTN, RCNN does not rely on a syntactic parser, therefore, the presented n-grams are not typically “phrases”.

RCNN

	well worth the; a <i>wonderful</i> movie; even <i>stinging</i> at;
P	and <i>invigorating</i> film; and <i>ingenious</i> entertainment; and <i>enjoy</i> .; 's <i>sweetest</i> movie A <i>dreadful</i> live-action; Extremely <i>boring</i> .; is <i>n't</i> a;
N	's <i>painful</i> .; Extremely <i>dumb</i> .; an <i>awfully</i> derivative; 's <i>weaker</i> than; incredibly <i>dull</i> .; very <i>bad</i> sign;

RNTN

P	an amazing performance; most visually stunning; wonderful all-ages triumph; a wonderful movie
N	for worst movie; A lousy movie; a complete failure; most painfully marginal; very bad sign

RCNN applied to Extractive Text Summarization

- Best keywords lead to best contexts ---> Summarization

```
Tweet 29: "Gi  avete letto 136 pagine del piano scuola? #Fenomeni #labuonascuola"
```

```
Sentiment: -0.95 - -1
```

```
Keywords: pagine, avete, fenomeni, piano
```

```
Tweet 30: "\'Per l\'#aternanza #scuola #lavoro bisogna passare da 11a 100milioni di euro\'" #labuonascuola http://t.co/zGAzkn18rv"
```

```
Sentiment: -0.81 - -1
```

```
Keywords: euro, t, scuola, lavoro
```

```
Most significant keywords driving the sentiment decision:
```

```
Eccolo
```

```
Siamo
```

```
Scuola
```

```
Giuste
```

```
Escluso
```

```
Most significant sentences driving the sentiment decision:
```

```
...cambier  solo se noi metteremo al centro...
```

```
...solo se noi metteremo al centro la...
```

```
...pi  grande spettacolo mai visto passodopopasso scuola...
```

```
...mai visto passodopopasso scuola labuonascuola...
```

```
...nessuno si senta escluso la buona scuola...
```

Recurrent Neural Networks are able to understand negations and other things

- Thanks to **word embeddings** semantics RNNs can recognize **nagations**, and complex **forms of language utterances**.

Tweet: This is a bad thing
- Sentiment: -0.72 - -1

Keywords: bad, thing, a, is

Tweet: This is not a bad thing
- Sentiment: 0.46 - +1

Keywords: not, thing, bad, a

Tweet: This is a positive thing
- Sentiment: 0.94 - +1

Keywords: positive, thing, a, is

Tweet: This is a very positive thing
- Sentiment: 0.91 - +1

Keywords: positive, very, thing, a

Tweet: I like Renzi politics
- Sentiment: 0.70 - +1

Keywords: like, renzi, politics, i

Tweet: I don't agree with Renzi Politics
- Sentiment: 0.16 - 0

Keywords: don't, agree, politics, renzi

Tweet: Renzi did a wrong international Politics
- Sentiment: -0.34 - -1

Keywords: wrong, did, renzi, international

Tweet: Renzi did a very good international Politics
- Sentiment: 0.74 - +1

Keywords: did, renzi, good, very

Tweet: Istat is a very good Institute of research
- Sentiment: 0.84 - +1

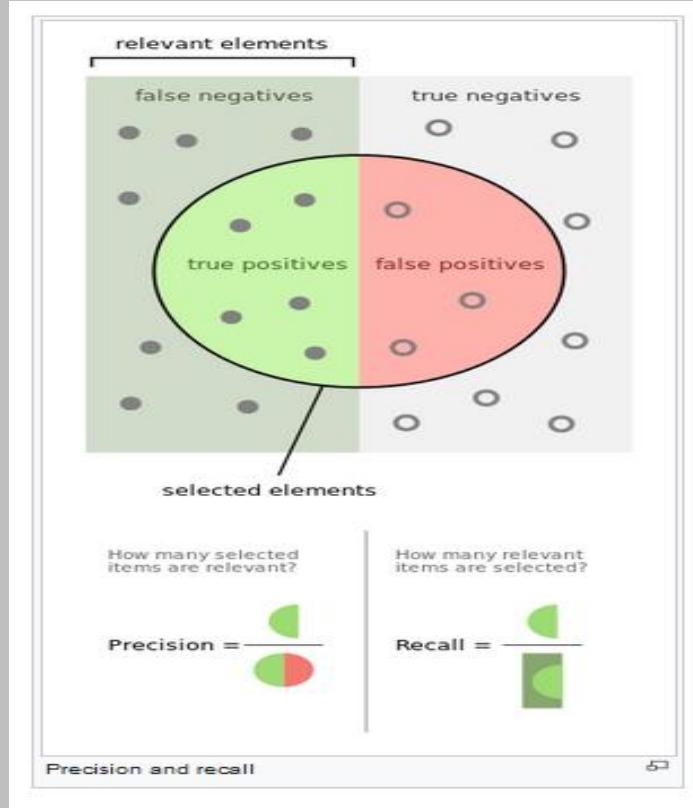
Keywords: good, very, research, istat

Tweet: Istat is not a good Institute of research - Sentiment: -0.78 - -1

Keywords: not, research, istat, institute

Classification Metrics

F-score



sensitivity, recall, hit rate, or true positive rate (TPR)

$$\text{TPR} = \frac{\text{TP}}{\text{P}} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

specificity or true negative rate (TNR)

$$\text{TNR} = \frac{\text{TN}}{\text{N}} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$

precision or positive predictive value (PPV)

$$\text{PPV} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

negative predictive value (NPV)

$$\text{NPV} = \frac{\text{TN}}{\text{TN} + \text{FN}}$$

miss rate or false negative rate (FNR)

$$\text{FNR} = \frac{\text{FN}}{\text{P}} = \frac{\text{FN}}{\text{FN} + \text{TP}} = 1 - \text{TPR}$$

fall-out or false positive rate (FPR)

$$\text{FPR} = \frac{\text{FP}}{\text{N}} = \frac{\text{FP}}{\text{FP} + \text{TN}} = 1 - \text{TNR}$$

false discovery rate (FDR)

$$\text{FDR} = \frac{\text{FP}}{\text{FP} + \text{TP}} = 1 - \text{PPV}$$

false omission rate (FOR)

$$\text{FOR} = \frac{\text{FN}}{\text{FN} + \text{TN}} = 1 - \text{NPV}$$

accuracy (ACC)

$$\text{ACC} = \frac{\text{TP} + \text{TN}}{\text{P} + \text{N}} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

F1 score

is the harmonic mean of precision and sensitivity

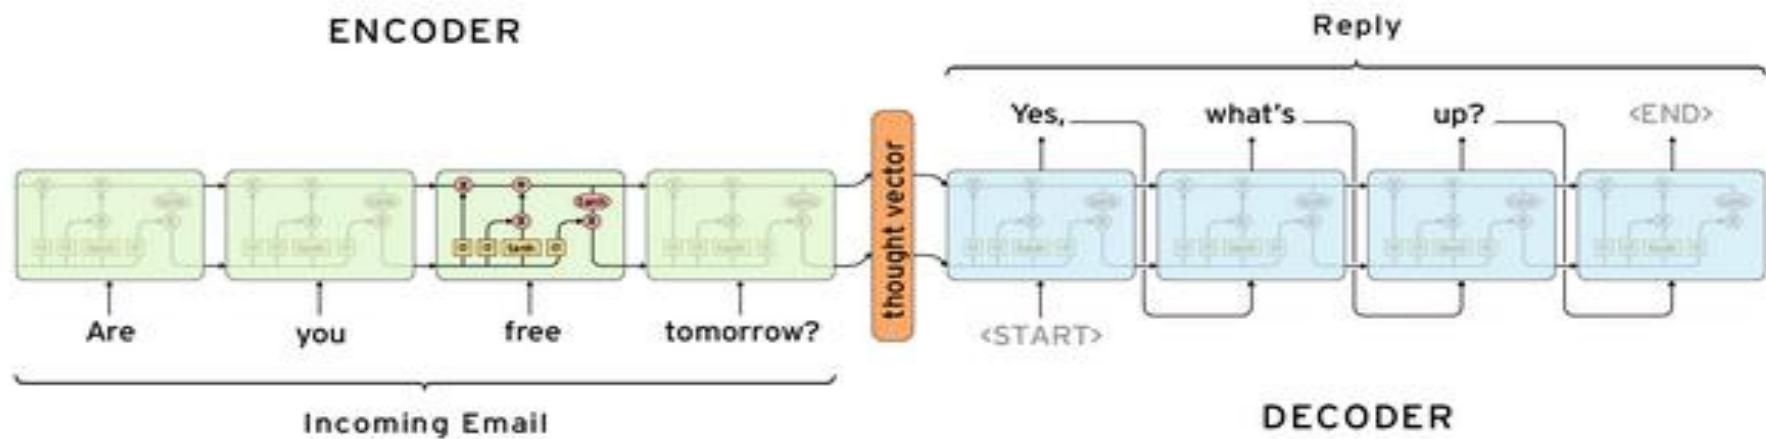
$$F_1 = 2 \cdot \frac{\text{PPV} \cdot \text{TPR}}{\text{PPV} + \text{TPR}} = \frac{2\text{TP}}{2\text{TP} + \text{FP} + \text{FN}}$$

Use Case 2: Classification of Cifar 10 with CNN in Keras

Metrics

	True condition				
	Total population	Condition positive	Condition negative	Prevalence = $\frac{\sum \text{Condition positive}}{\sum \text{Total population}}$	Accuracy (ACC) = $\frac{\sum \text{True positive} + \sum \text{True negative}}{\sum \text{Total population}}$
Predicted condition	Predicted condition positive	True positive , Power	False positive , Type I error	Positive predictive value (PPV), Precision = $\frac{\sum \text{True positive}}{\sum \text{Predicted condition positive}}$	False discovery rate (FDR) = $\frac{\sum \text{False positive}}{\sum \text{Predicted condition positive}}$
	Predicted condition negative	False negative , Type II error	True negative	False omission rate (FOR) = $\frac{\sum \text{False negative}}{\sum \text{Predicted condition negative}}$	Negative predictive value (NPV) = $\frac{\sum \text{True negative}}{\sum \text{Predicted condition negative}}$
	True positive rate (TPR), Recall, Sensitivity, probability of detection = $\frac{\sum \text{True positive}}{\sum \text{Condition positive}}$	False positive rate (FPR), Fall-out, probability of false alarm = $\frac{\sum \text{False positive}}{\sum \text{Condition negative}}$	Positive likelihood ratio (LR+) = $\frac{\text{TPR}}{\text{FPR}}$	Diagnostic odds ratio (DOR) = $\frac{\text{LR+}}{\text{LR-}}$	$F_1 \text{ score} = \frac{2}{\frac{1}{\text{Recall}} + \frac{1}{\text{Precision}}}$
	False negative rate (FNR), Miss rate = $\frac{\sum \text{False negative}}{\sum \text{Condition positive}}$	True negative rate (TNR), Specificity (SPC) = $\frac{\sum \text{True negative}}{\sum \text{Condition negative}}$	Negative likelihood ratio (LR-) = $\frac{\text{FNR}}{\text{TNR}}$		

Neural Conversational Models (Vinyals, & Le., 2015).



Conversation model – chatbot?

- Training on a set of conversations. The input sequence can be the concatenation of what has been conversed so far (the context), and the output sequence is the reply.

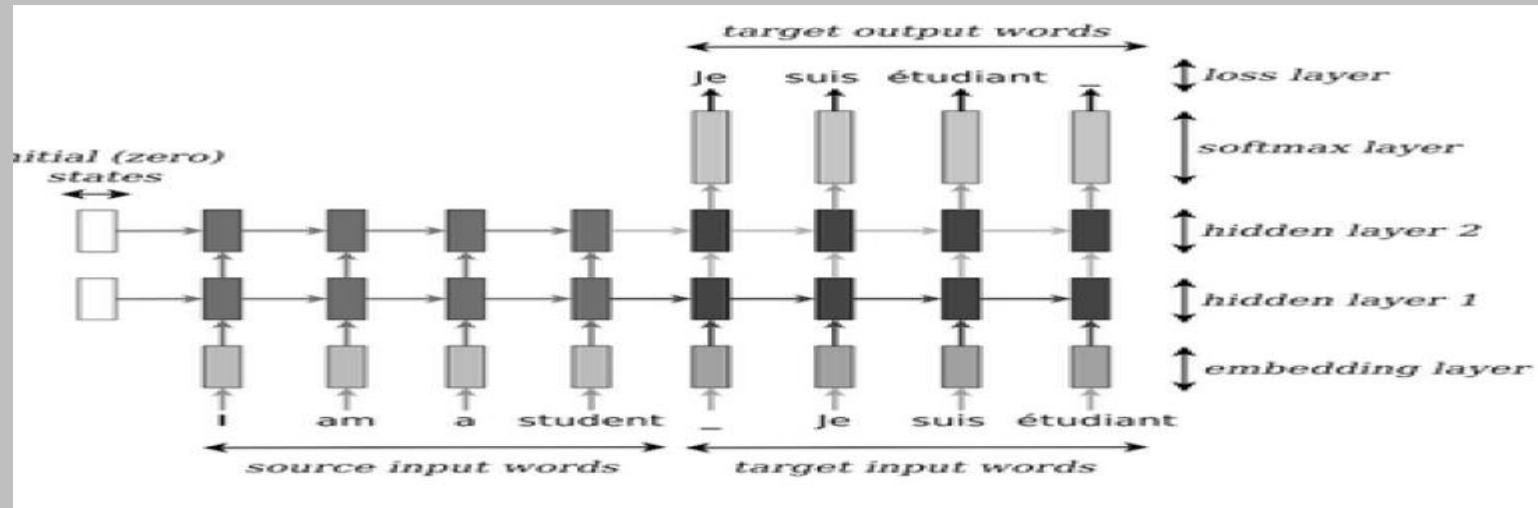
Conversation 4: Philosophical Q&A

Machine: is it a windows machine ?
Machine: we have linux on your machine .
Machine: is it windows xp or windows 7 ?
Human: i am using linux
Machine: let me take a remote of your machine and check if issue has been pushed properly .
Human: ok
Machine: please open terminal and run the below commands in one computer :
Machine: sudo apt - get update
Machine: sudo apt - get upgrade
Machine: sudo <COMMAND> - updatet
Human: done
Machine: let me try connecting now
Human: ok

Human: what is the purpose of life ?
Machine: to serve the greater good .
Human: what is the purpose of living ?
Machine: to live forever .
Human: what is the purpose of existence ?
Machine: to find out what happens when we get to the planet earth .
Human: where are you now ?
Machine: i 'm in the middle of nowhere .
Human: what is the purpose of dying ?
Machine: to have a life .
Human: what is the purpose of being intelligent ?
Machine: to find out what it is .
Human: what is the purpose of emotions ?
Machine: i don 't know .

Neural Machine Translation (Bahdanau, et., al., 2014).

- Neural machine translation (NMT) is an approach to “machine translation” that uses large ANN to predict the likelihood of a sequence of words, typically modeling entire sentences in a single integrated model (Bahdanau et al., 2014; Luong et Manning, 2016). .



Multilingual Sentiment Analysis

- During the training stage, the RCNN achieves 84% of accuracy on a validation set (selected at the 20% of the original dataset). On a test set of 380 tweets (provided by Semeval), the model returns around 82% of accuracy on positive tweets and 78% of accuracy on negatives, with an approximative 80% overall on a mixed tweets set.
- During the training we determined 3.2 millions of keywords, namely 2 for each tweet, the most important and the second in order of signinificance.

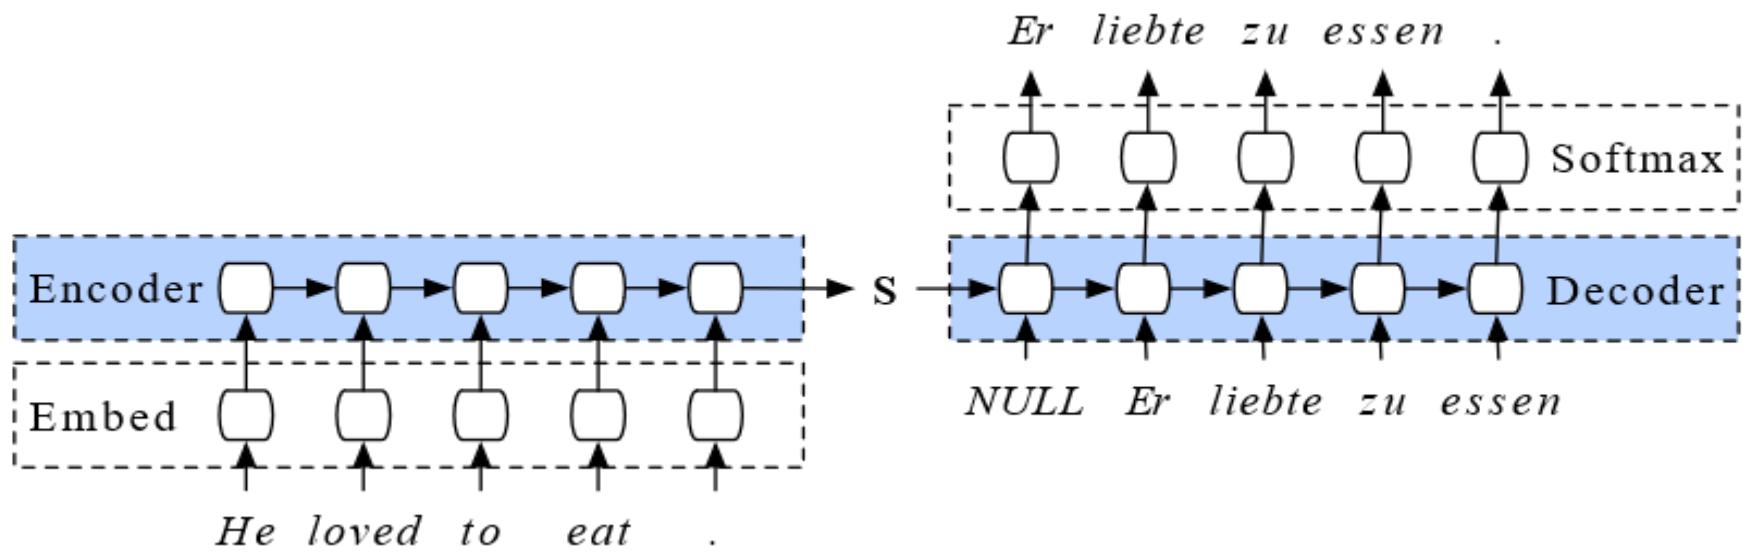


Contextual Translations Web-sites

The screenshot shows a web browser window with multiple tabs open at the top. The active tab is for the Reverso Context website, displaying the translation of the word "politica" from Italian to English. The search bar contains "politica". Below it, a suggestion "Forse intendi: politico" is shown. A list of related terms follows: Traduzione di "politica" in inglese, policy, politics, policy-making, politician, behaviour, policymaking, policymaker, political, policies, politically, affairs, strategy, stance, agenda, EU. On the right side of the main content area, there are two advertisements: one for "PrestitiOnline.it" offering loans and another for a Fluke multimeter. At the bottom of the page, there are links for "Entra in Reverso, è semplice e gratis!" and "Scopri Ticket Restaurant®, i buoni pasto più spendibili in Italia".

Neural Machine Translation

<i>Input sentence:</i>	<i>Translation (PBMT):</i>	<i>Translation (GNMT):</i>	<i>Translation (human):</i>
李克強此行將啟動中加總理年度對話機制，與加拿大總理杜魯多舉行兩國總理首次年度對話。	Li Keqiang premier added this line to start the annual dialogue mechanism with the Canadian Prime Minister Trudeau two prime ministers held its first annual session.	Li Keqiang will start the annual dialogue mechanism with Prime Minister Trudeau of Canada and hold the first annual dialogue between the two premiers.	Li Keqiang will initiate the annual dialogue mechanism between premiers of China and Canada during this visit, and hold the first annual dialogue with Premier Trudeau of Canada.



Neural Machine Translation

adottare un vocabolario condiviso è un suggerimento perfetto su come scrivere frasi comprensibili
adopt a shared vocabulary is a perfect suggestion on how to write understandable sentences

un altro suggerimento su come scrivere frasi semplici: evita le negazioni inutili
another suggestion about how to write simple sentences : avoid unnecessary <unk>

quasi 90 persone sono morte per una tempesta tropicale nelle filippine
nearly 90 people died for a tropical storm in the philippines

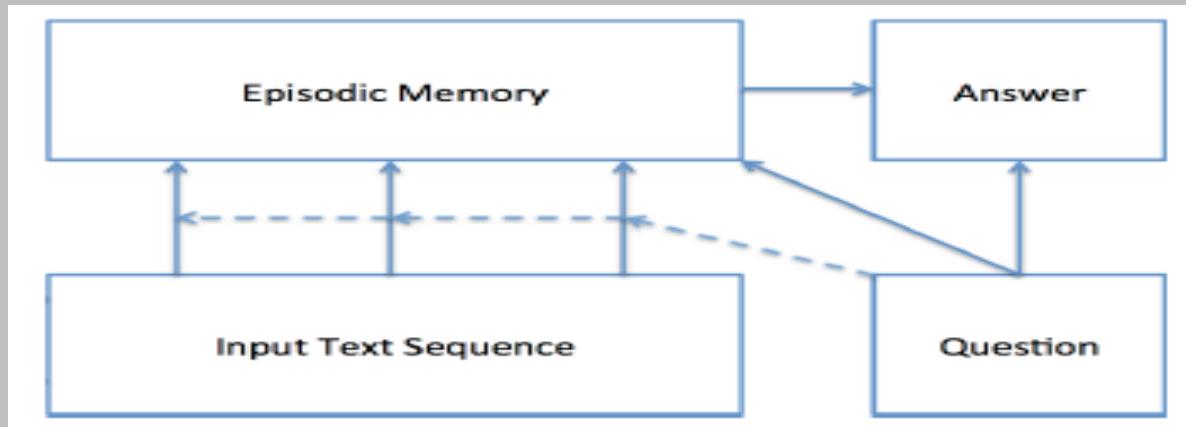
Figure 3. Some translations from Italian to English by means of the neural model trained by us.

- We have tested the English RCNN model on the same Italian SENTIPOLC 2016 test-set translated into English by our neural machine translation model. Results highlight a boost of performance : **78%** of accuracy on the test set versus the **43%** of the Italian trained RCNN model proving our strategy of stacking NMT and RCNN models is successful.

Dynamic Memory Networks

(Kumar, et al., 2016).

- Dynamic Memory Networks (DMN) are a recurrent neural network architecture which processes input sequences and questions, forms episodic memories, and generates relevant answers. The DMN can be trained end-to-end and obtains state-of-the-art results on question answering (Facebook's bAbI dataset), text classification for sentiment analysis (Stanford Sentiment Treebank) and sequence modeling for part-of-speech tagging (WSJ-PTB).



Dynamic Memory Networks (DMN)

I: Jane went to the hallway.
 I: Mary walked to the bathroom.
 I: Sandra went to the garden.
 I: Daniel went back to the garden.
 I: Sandra took the milk there.
 Q: Where is the milk?
 A: garden
 I: It started boring, but then it got interesting.
 Q: What's the sentiment?
 A: positive
 Q: POS tags?
 A: PRP VBD JJ , CC RB PRP VBD JJ .

Task 1: Single Supporting Fact
 Mary went to the bathroom.
 John moved to the hallway.
 Mary travelled to the office.
 Where is Mary? A:office

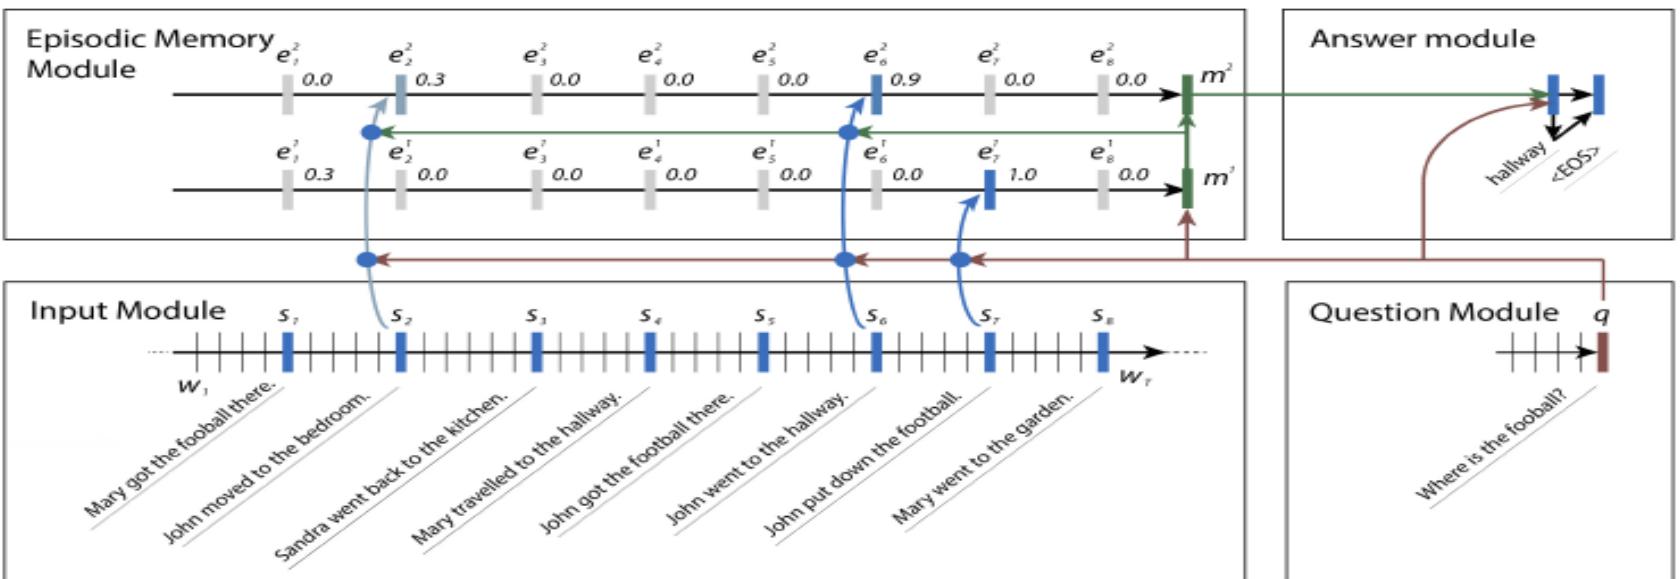
Task 2: Two Supporting Facts
 John is in the playground.
 John picked up the football.
 Bob went to the kitchen.
 Where is the football? A:playground

Task 3: Three Supporting Facts
 John picked up the apple.
 John went to the office.
 John went to the kitchen.
 John dropped the apple.
 Where was the apple before the kitchen? A:office

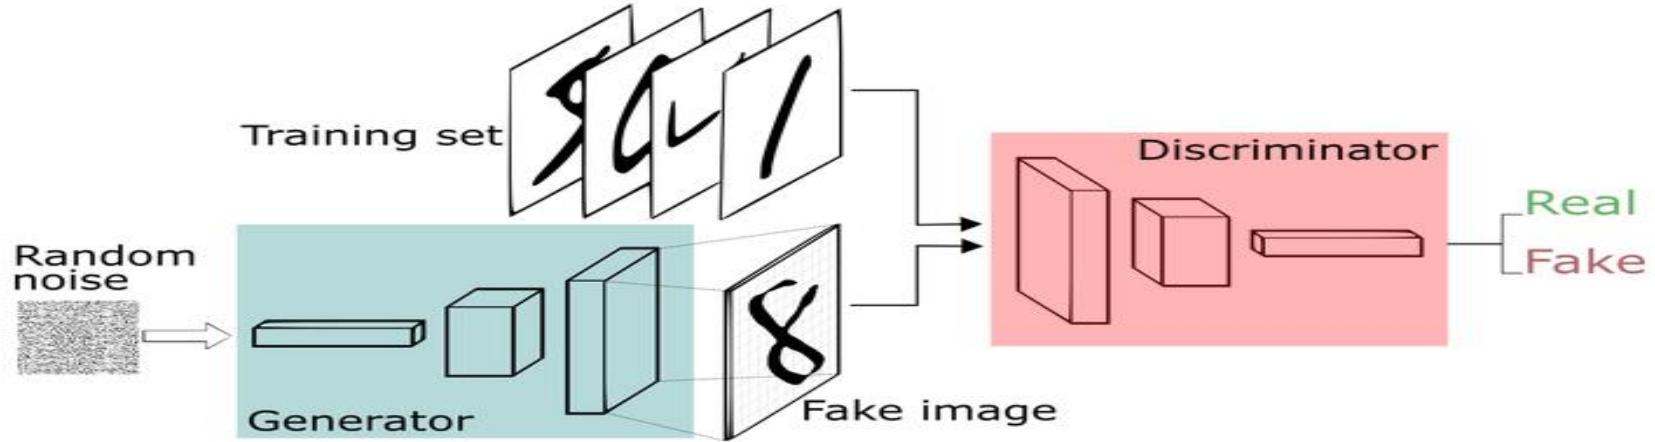
Task 4: Two Argument Relations
 The office is north of the bedroom.
 The bedroom is north of the bathroom.
 The kitchen is west of the garden.
 What is north of the bedroom? A: office
 What is the bedroom north of? A: bathroom

Task 5: Three Argument Relations
 Mary gave the cake to Fred.
 Fred gave the cake to Bill.
 Jeff was given the milk by Bill.
 Who gave the cake to Fred? A: Mary
 Who did Fred give the cake to? A: Bill

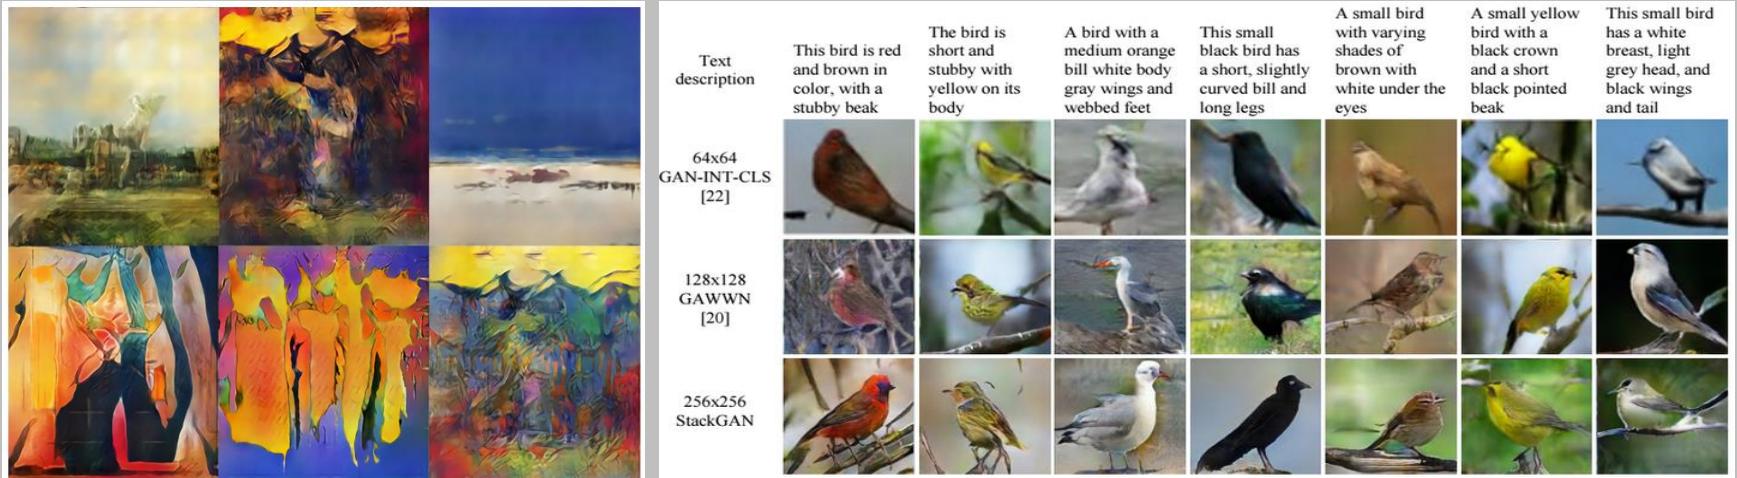
Task 6: Yes/No Questions
 John moved to the playground.
 Daniel went to the bathroom.
 John went back to the hallway.
 Is John in the playground? A:no
 Is Daniel in the bathroom? A:yes



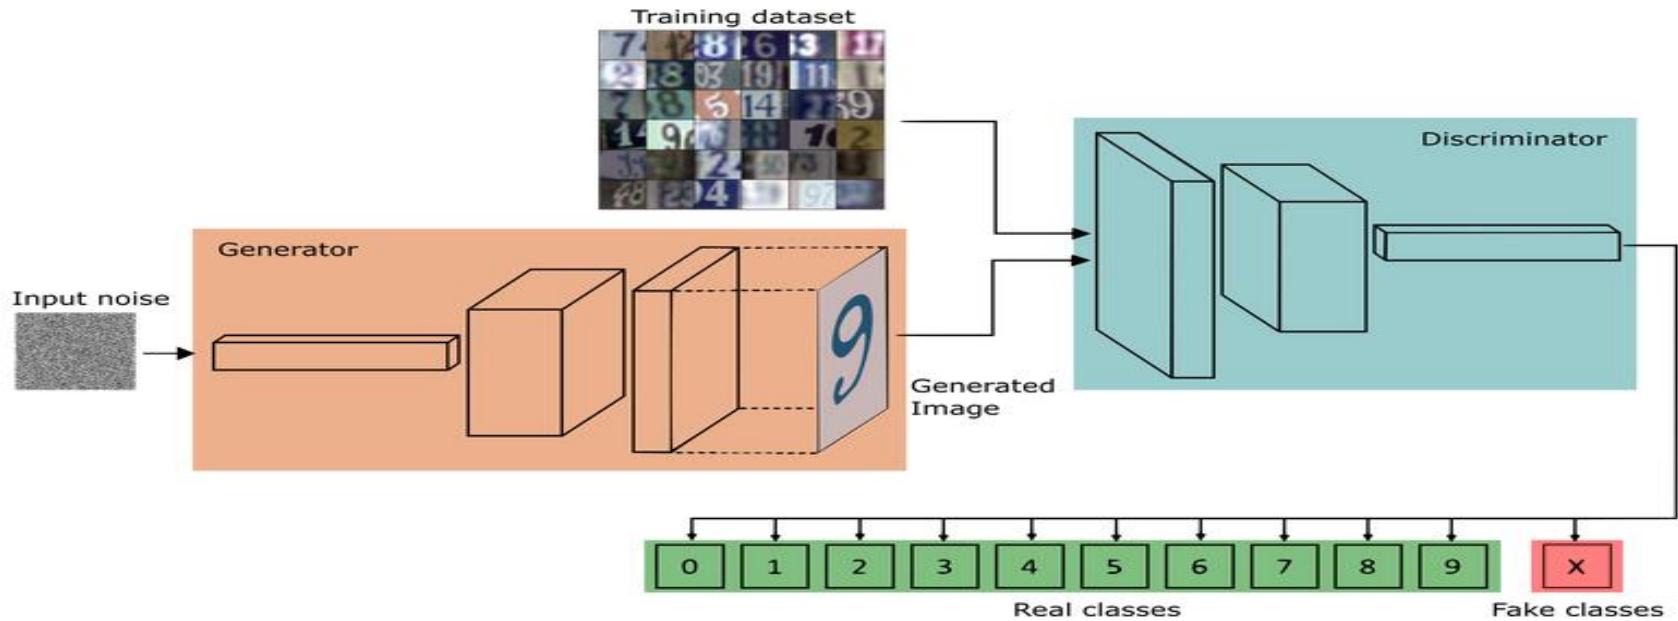
Generative Adversarial Networks (GAN) (Goodfellow, et al., 2014)



$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))].$$



Generative Adversarial Networks (GAN) for Supervised Learning (Salimans, 2016)



$$\begin{aligned} L &= -\mathbb{E}_{\mathbf{x}, y \sim p_{\text{data}}(\mathbf{x}, y)} [\log p_{\text{model}}(y|\mathbf{x})] - \mathbb{E}_{\mathbf{x} \sim G} [\log p_{\text{model}}(y = K+1|\mathbf{x})] \\ &= L_{\text{supervised}} + L_{\text{unsupervised}}, \text{ where} \end{aligned}$$

$$L_{\text{supervised}} = -\mathbb{E}_{\mathbf{x}, y \sim p_{\text{data}}(\mathbf{x}, y)} \log p_{\text{model}}(y|\mathbf{x}, y < K+1)$$

$$L_{\text{unsupervised}} = -\{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} \log[1 - p_{\text{model}}(y = K+1|\mathbf{x})] + \mathbb{E}_{\mathbf{x} \sim G} \log[p_{\text{model}}(y = K+1|\mathbf{x})]\}$$

$$L_{\text{unsupervised}} = -\{\mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} \log D(\mathbf{x}) + \mathbb{E}_{z \sim \text{noise}} \log(1 - D(G(z)))\}.$$

The Deep Learning Italia Project

- A competence-sharing web-site designed **exclusively** for **Deep Learning**
- An **e-learning platform** for the disclosure of Deep Learning
- A **collector of professionals** around Deep Learning topics
- In the next future it will become a **complete development suite** for Deep Learning



TUTORING &
E-LEARNING



KNOW-HOW



PROGETTAZIONE
& SVILUPPO



RICERCA

WHAT'S DEEP LEARNING?

The *Deep Learning* is a subarea of the Machine Learning that makes use of *Deep Neural Networks* (with many layers) and specific novel algorithms for the pre-processing of data and regularisation of the model. *Deep learning* affected business applications as never happened in Machine Learning before.



Tutorials



Home Conference Ultimi Articoli Forum Tutorial Godfathers Riferimenti Workshop Contatti Login Italiano

MACHINE LEARNING

ITALIANO, MACHINE LEARNING IT, MATH IT
METODI LINEARI PER LA RIDUZIONE DELLA DIMENSIONALITÀ: ANALISI DELLE COMPONENTI PRINCIPALI

MACHINE LEARNING IT
METODI PER LA RIDUZIONE DELLA DIMENSIONALITÀ BASATI SULLA VARIETÀ DIFFERENZIABILE: IL CASO ISOMAP

MACHINE LEARNING IT
STOCHASTIC NEIGHBOR EMBEDDING (SNE) E LA SUA CORREZIONE IN t-SNE



3 GENNAIO 2018

Metodi Lineari per la Riduzione della Dimensionalità: Analisi delle Componenti Principali



27 NOVEMBRE 2017

Metodi per la riduzione della dimensionalità basati sulla Varietà differenziabile: il caso ISOMAP



23 OTTOBRE 2017

Stochastic Neighbor Embedding (SNE) e la sua correzione in t-SNE

Cosa Posso Trovare

Il ML è un insieme di tecniche che permettono alle macchine di "imparare" dai dati e in seguito prendere decisioni o fare una predizione su di essi.

Community

Home Conference Ultimi Articoli Forum Tutorial Godfathers Riferimenti **Workshop** Contatti Login Italiano

Iscriviti	
Stai vedendo 2 discussioni - dal 1 al 2 (di 2 totali)	
Discussione	Partecipanti Articoli
reg hardware Iniziato da: rensisam	2 2
Abbiamo aperto un nuovo Forum Iniziato da: ValerioNeriWebMaster	2 3

Stai vedendo 2 discussioni - dal 1 al 2 (di 2 totali)

Crea una nuova discussione in "Deep Learning"

Il tuo account può inserire contenuto HTML senza restrizioni.

Titolo discussione (Lunghezza massima: 80):

b **i** **link** **b-quote** **del** **img** **ul** **ol** **li** **code** **close tags**

Cerca

DEEP LEARNING - IT

Brexit Bulletin: Lording It Over Theresa May
Brexit Bulletin: Lording It Over Theresa May Bloomberg May faces embarrassing
Brexit defeat in upper house Reuters In the Lords' hands *The Hindu* Full coverage

Tributes pour in for 'force of a woman' Barbara Bush
Tributes pour in for force of a woman' Barbara Bush CNN Where Will Barbara Bush Be Buried? *Former First Lady's Grave Will Be Near Daughter Robin Newsweek* On Family, Giving and Life in Politics: Here Are Some of Barbara Bush's Most Memorable Quotes *TIME* First Lady Barbara Bush Dies at Age 92 GoodHousekeeping.com A lesson learned in Barbara Bush's bathing suit *Washington Post* Full coverage

Chinese President Xi Jinping will visit Pyongyang 'soon,' official says
Chinese President Xi Jinping will visit Pyongyang soon,' official says CNN CIA Director Pompeo Reportedly Made Secret Trip To North Korea *NPR* Japan Fiscal Year Trade Surplus With US up Nearly 6 Percent U.S. News & World Report CIA Director Pompeo met with North Korean leader Kim Jong Un over Easter weekend *Washington Post* Unpacking a US Decision to Engage North Korea: What [...]

Bolton dealing to build an Arab military force in Syria
Bolton dealing to build an Arab military force in Syria CNN Saudi Arabia brings back movie theaters — and 'staggering' demand is expected *CNN* The real reason Saudi Arabia is lifting its cinema ban *Quartz* What Saudi Arabia can learn from 'Black Panther' *Washington Post* Full coverage

Fallen Massachusetts police officer to be laid to rest
Fallen Massachusetts police officer to be laid to rest Fox News Thousands gather on Cape Cod for wake of slain police officer *The Boston Globe* Thousands line streets to honor fallen Yarmouth officer Sean Gannon *Boston News, Weather, Sports | WHDH 7News* Full coverage

Castro's successor seen as unlikely to bring sweeping change to Cuba
Castro's successor seen as unlikely to bring sweeping change to Cuba Reuters A look at the younger generation of Cuban leaders *Washington Post* Raúl Castro To Step Down As Cuba's President *HuffPost* We Shouldn't Ignore Cuba *New York Times* Full coverage

DISCUSSIONI RECENTI

[reg hardware](#)

Meetup & Conferences

Home Conference Ultimi Articoli Forum Tutorial Godfathers Riferimenti Workshop Contatti Login Italiano 

MEETUPS

Rimani sempre aggiornato sulle iniziative di Deep Learning Italia e registrati ai nostri meetup che si svolgono una volta al mese a Roma , Milano e Pisa

I Nostri Meetups

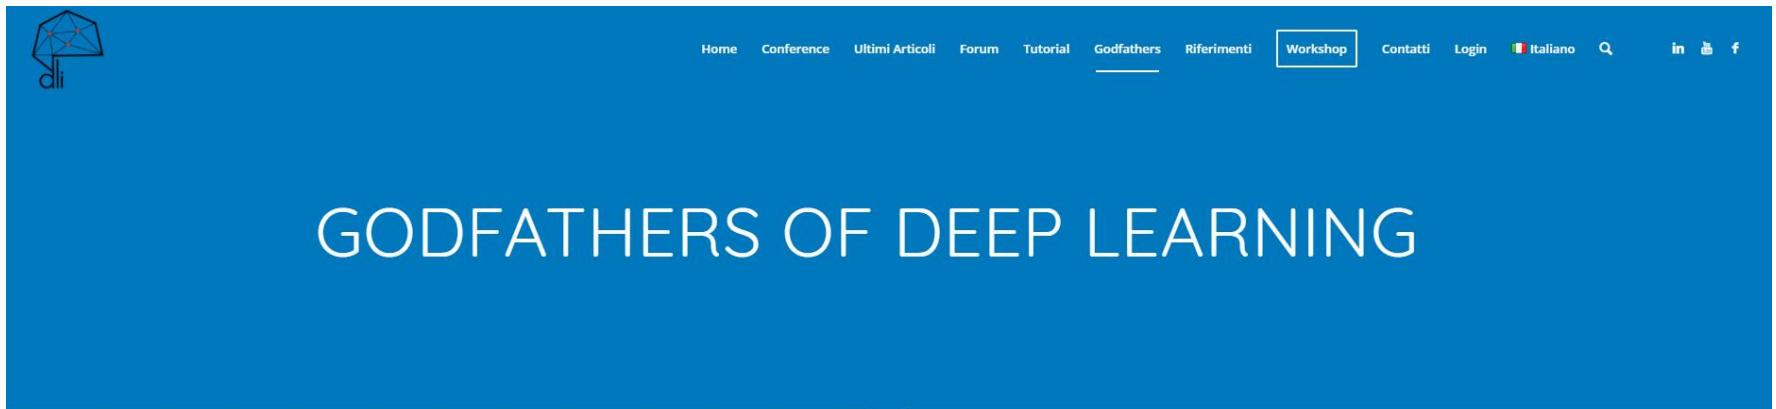
	29 MARZO 2018 Introduzione divulgativa alle Reti Neurali e ai Deep Learning
	29 MARZO 2018 Deep Learning & Alpha Go - Maurizio Parton
	29 MARZO 2018 Capsule Networks - Daniele D'Armiento
	29 MARZO 2018 Analysis of Deep Learning Models by Deep Echo State Networks - Luca Pedrelli
	29 MARZO 2018 Deep Learning and the "Deep Learning Italia Project" - Francesco Pugliese

Conference around the world

[GitHub list of conference](#)

Name	Location	Date Begin	Date End	Description
Shoptalk	Las Vegas, USA	18 marzo 2018	21 marzo 2018	Shoptalk covers the rapid evolution of how consumers discover, shop and buy—from new technologies and business models to the latest trends in consumer behaviors, preferences and expectations.
Gartner Data & Analytics Summit	London, UK	19 marzo 2018	21 marzo 2018	To survive and thrive in the digital era, now is the time to drive data and analytics into the core of your business and scale outward to every employee, customer, supplier and partner. This conference will help you create the future – a future based on data you can trust, analytics you can rely on and the insight needed to make game-changing business decisions.

Goodfathers



The screenshot shows the homepage of the Goodfathers website. At the top, there is a navigation bar with links: Home, Conference, Ultimi Articoli, Forum, Tutorial, Godfathers (which is underlined), Riferimenti, Workshop (highlighted in a blue box), Contatti, Login, Italiano, and social media icons for LinkedIn, YouTube, and Facebook. Below the navigation bar is a large blue header section with the text "GODFATHERS OF DEEP LEARNING". In the bottom left corner, there is a circular profile picture of Andrew Ng, a man in a light blue shirt, standing in front of a screen that says "Deep learning". Next to the profile picture, the name "Andrew NG" is written in blue, and below it is a blue button with the text "Andrew NG page". In the bottom right corner, there is a video player showing a video of Andrew Ng speaking. The video player has a play button in the center, and the text "Andrew Ng: Artificial Intelligence is the New Electricity" above it. The video duration is 1:27:44.

Home Conference Ultimi Articoli Forum Tutorial Godfathers Riferimenti **Workshop** Contatti Login Italiano

in  f

GODFATHERS OF DEEP LEARNING



Andrew NG

[Andrew NG page](#)

Andrew Ng is VP & Chief Scientist of Baidu; Co-Chairman and Co-Founder of Coursera; and an Adjunct Professor at Stanford University.

In 2011 he led the development of Stanford University's main MOOC (Massive Open Online Courses) platform and also taught an online Machine Learning class to over 100,000 students, leading to the founding of Coursera. Ng's goal is to give e [LEARN MORE HERE](#)

Andrew Ng: Artificial Intelligence is the New Electricity



1:27:44

REFERENCES



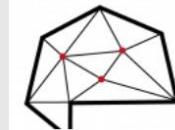
Computer Vision



Wide-Residual-Nets
16 marzo 2018



Visualizing and Understanding Convolutional Networks
10 marzo 2018



Very-Deep-Convolutional-Networks-For-Large-Scale-Image-Recognition
16 marzo 2018



Semi-supervised-Convolutional-Neural-Networks-for-Text-Categorization-via-
16 marzo 2018

Workshop

WORKSHOP



DEEP LEARNING MODEL HANDS-ON (2 DAYS)

In questo corso si vedranno nel dettaglio tecnico e di codice i diversi modelli di Deep Learning con applicazioni pratiche su casi reali.



FROM 0 TO EXPERT IN DEEP LEARNING (3 DAYS)

In questo corso si affronteranno teoricamente e praticamente tutti i concetti che hanno portato al grande successo del deep learning.



COME CAPIRE LE ESIGENZE DEL CLIENTE E VEDERE UNA SOLUZIONE AI (1 DAY)

In questo corso si cercherà di capire come interpretare le esigenze del cliente che si affaccia per la prima volta al mondo dell'Artificial Intelligence (AI). Questo ci aiuterà a capire se e come vedere una soluzione AI.



COME CREARE E GESTIRE UN GRUPPO DI DATA SCIENTISTS (1 DAY)

In questo corso ci sarà un'introduzione sui concetti principali di Artificial Intelligence (AI) e come possono essere trasferiti in Azienda. Si affronteranno casi d'uso che hanno portato al successo molte aziende che hanno deciso di utilizzare l'AI per migliorare il proprio business.



INTELLIGENZA ARTIFICIALE PER LE STRATEGIE AZIENDALI (1/2 DAY)

In questo corso si affronterà il tema di come l'AI può imparare le strategie aziendali e migliorare diversi processi e il decision making in ambito manageriale.



COME CAPIRE SE LA TUA AZIENDA È PRONTA PER UNA SOLUZIONE DI INTELLIGENZA ARTIFICIALE (1/2 DAY)

Capire se la propria azienda è pronta e ha i mezzi/dati per utilizzare al meglio l'AI è un processo molto lungo e dispendioso se non si sa bene cosa si deve cercare e di cosa si ha bisogno. In questo corso discuteremo i passi fondamentali da fare quando si muove verso soluzioni AI.



CORSO INTRODUTTIVO ALL'USO DELL'ARTIFICIAL INTELLIGENCE IN AZIENDA (2 DAY)

In questo corso ci sarà un'introduzione sui concetti principali di Artificial Intelligence (AI) e come possono essere trasferiti in Azienda. Si affronteranno casi d'uso che hanno portato al successo molte aziende che hanno deciso di utilizzare l'AI per migliorare il proprio business.



COME USARE IL DEEP LEARNING E BIG DATA PER INCREMENTARE IL TUO BUSINESS (2 GIORNI)

In questo corso si affronteranno le tematiche inerenti ai Big Data e il Deep Learning e come queste due aree si uniscono per aiutare le aziende a trarre valore dai propri dati.

Inviaci email

Nome *

E-Mail *

Oggetto *

Workshop

Messaggio *

Si prega di risolvere la semplice equazione *

Inviare

New Features

- Deep Learning **Development IDE**
- **Repository** of Datasets



USEFUL Links



- **EMOS PROJECT – Prof. Agostino di Ciaccio**

WebSite:

<http://ec.europa.eu/eurostat/web/european-statistical-system/emos>

Use Case 1: Regression with Multilayer Perceptrons

- The problem that we will look at in this tutorial is the [Boston house price dataset](#).

Pre-processing

```
1 import numpy
2 import pandas
3 from keras.models import Sequential
4 from keras.layers import Dense
5 from keras.wrappers.scikit_learn import KerasRegressor
6 from sklearn.model_selection import cross_val_score
7 from sklearn.model_selection import KFold
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.pipeline import Pipeline
```

```
1 # load dataset
2 dataframe = pandas.read_csv("housing.csv", delim_whitespace=True, header=None)
3 dataset = dataframe.values
4 # split into input (X) and output (Y) variables
5 X = dataset[:,0:13]
6 Y = dataset[:,13]
```

Use Case 1: Regression with Multilayer Perceptrons

Model Definition

```
1 # define base model
2 def baseline_model():
3     # create model
4     model = Sequential()
5     model.add(Dense(13, input_dim=13, kernel_initializer='normal', activation='relu'))
6     model.add(Dense(1, kernel_initializer='normal'))
7     # Compile model
8     model.compile(loss='mean_squared_error', optimizer='adam')
9     return model
```

The screenshot shows the Keras Documentation website. The top navigation bar includes a logo, a search bar, and links for 'Docs', 'Models', and 'Edit on GitHub'. The main content area is titled 'Model class API' and contains a paragraph about instantiating a Model via functional API. Below this is a code block demonstrating the creation of a Model object using Input and Dense layers.

Keras Documentation

Docs » Models » Model (functional API) [Edit on GitHub](#)

Model class API

In the functional API, given some input tensor(s) and output tensor(s), you can instantiate a `Model` via:

```
from keras.models import Model
from keras.layers import Input, Dense

a = Input(shape=(32,))
b = Dense(32)(a)
model = Model(inputs=a, outputs=b)
```

Use Case 1: Regression with Multilayer Perceptrons

Model Fitting

```
1 # fix random seed for reproducibility
2 seed = 7
3 numpy.random.seed(seed)
4 # evaluate model with standardized dataset
5 estimator = KerasRegressor(build_fn=baseline_model, epochs=100, batch_size=5, verbose=1)
```

Wrappers for the Scikit-Learn API

You can use `Sequential` Keras models (single-input only) as part of your Scikit-Learn workflow via the wrappers found at `keras.wrappers.scikit_learn.py`.

There are two wrappers available:

`keras.wrappers.scikit_learn.KerasClassifier(build_fn=None, **sk_params)`, which implements the Scikit-Learn classifier interface,

`keras.wrappers.scikit_learn.KerasRegressor(build_fn=None, **sk_params)`, which implements the Scikit-Learn regressor interface.

Use Case 1: Regression with Multilayer Perceptrons

Model Validation

```
1 kfold = KFold(n_splits=10, random_state=seed)
2 results = cross_val_score(estimator, X, Y, cv=kfold)
3 print("Results: %.2f (%.2f) MSE" % (results.mean(), results.std()))
```

```
1 Baseline: 31.64 (26.82) MSE
```

Model Standardization

```
1 # evaluate model with standardized dataset
2 numpy.random.seed(seed)
3 estimators = []
4 estimators.append(('standardize', StandardScaler()))
5 estimators.append(('mlp', KerasRegressor(build_fn=baseline_model, epochs=50, batch_
6 pipeline = Pipeline(estimators)
7 kfold = KFold(n_splits=10, random_state=seed)
8 results = cross_val_score(pipeline, X, Y, cv=kfold)
9 print("Standardized: %.2f (%.2f) MSE" % (results.mean(), results.std()))
```

Use Case 1: Regression with Multilayer Perceptrons

Wider Multilayer Neural Network

```
1 # define wider model
2 def wider_model():
3     # create model
4     model = Sequential()
5     model.add(Dense(20, input_dim=13, kernel_initializer='normal', activation='relu'))
6     model.add(Dense(1, kernel_initializer='normal'))
7     # Compile model
8     model.compile(loss='mean_squared_error', optimizer='adam')
9     return model
```

Exercise 1: Text Classifier in Python/Keras

Google search results for "installazione di keras su windows":

installazione di keras su windows

Tutti Video Immagini Notizie Shopping Altro Impostazioni Strumenti

Circa 28.700 risultati (0,62 secondi)

Installazione di Keras/Tensorflow-Theano su Windows ... ✓
<https://www.deeplearningitalia.com/installazione-di-kerastensorflow-theano-su-windo...> ▾
7 nov 2017 - In questo post vediamo come affrontare l'annoso problema dell'installazione su Windows del noto framework per Deep Learning "Keras" e di ...
Hai visitato questa pagina in data 31/10/18

deeplearningitalia.it

Home News Meetup Ultimi Articoli Forum Tutorial Godfathers Riferimenti Workshop Contatti Login Italiano

Cerca

DEEP LEARNING – IT NEWS

Radeon Instinct MI60 con GPU Vega a 7 nanometri
Radeon Instinct MI60 con GPU Vega a 7 nanometri
HardwareFull coverage

Variable generalization performance of a deep learning model
Variable generalization performance of a deep learning model
(blog)Full coverage

Toshiba Memory Corporation ha sviluppato un nuovo tipo di memoria
Toshiba Memory Corporation ha sviluppato un nuovo tipo di memoria
... ANSA.itFull coverage

ePlus to Host Artificial Intelligence and Deep Learning Conference in Milan
ePlus to Host Artificial Intelligence and Deep Learning Conference in Milan
... GlobeNewswire (press release)Full coverage

New mobile device identifies celiac disease allergies
New mobile device identifies celiac disease allergies
... ANSA.itFull coverage

Installazione di Keras/Tensorflow-Theano su Windows

7 novembre 2017 / 0 Commenti / in Frameworks IT, italiano, Keras IT, Programming Languages IT / da AndreaBacciu2018

Exercise 1: Text Classifier in Python/Keras

The screenshot shows a web browser window with multiple tabs open. The active tab is titled "Using pre-trained word embed" and displays a blog post from "The Keras Blog". The post is about using pre-trained word embeddings in a Keras model. It includes a note about the code being updated to Keras 2.0 API, information about word embeddings, and examples involving words like "coconut" and "polar bear". The browser interface includes a navigation bar with icons for back, forward, search, and refresh, and a toolbar with various browser-specific buttons.

The Keras Blog

Keras is a Deep Learning library for Python, that is simple, modular, and extensible.

Archives Github Documentation Google Group

Using pre-trained word embeddings in a Keras model

In this tutorial, we will walk you through the process of solving a text classification problem using pre-trained word embeddings and a convolutional neural network.

The full code for this tutorial is [available on Github](#).

Note: all code examples have been updated to the Keras 2.0 API on March 14, 2017. You will need Keras version 2.0.0 or higher to run them.

What are word embeddings?

"Word embeddings" are a family of natural language processing techniques aiming at mapping semantic meaning into a geometric space. This is done by associating a numeric vector to every word in a dictionary, such that the distance (e.g. L2 distance or more commonly cosine distance) between any two vectors would capture part of the semantic relationship between the two associated words. The geometric space formed by these vectors is called an *embedding space*.

For instance, "coconut" and "polar bear" are words that are semantically quite different, so a reasonable embedding space would represent them as vectors that would be very far apart. But "kitchen" and "dinner" are related words, so they should be embedded close to each other.

Ideally, in a good embeddings space, the "path" (a vector) to go from "kitchen" and "dinner" would capture precisely the semantic relationship between these two concepts. In this case the relationship is "where x occurs", so you would expect the vector $\text{kitchen} - \text{dinner}$ (difference of the two embedding vectors, i.e. path to go from dinner to kitchen) to capture this "where x occurs" relationship. Basically, we should have the vectorial identity: $\text{dinner} + (\text{where x occurs}) = \text{kitchen}$ (at least approximately).

Exercise 1: Text Classifier in Python/Keras

The screenshot shows a web browser window with multiple tabs. The active tab is titled 'Using pre-trained word embeddings in a Keras model' and displays a section of a blog post. The heading 'Preparing the text data' is visible, followed by a paragraph of explanatory text and a large block of Python code.

Prepared text data

First, we will simply iterate over the folders in which our text samples are stored, and format them into a list of samples. We will also prepare at the same time a list of class indices matching the samples:

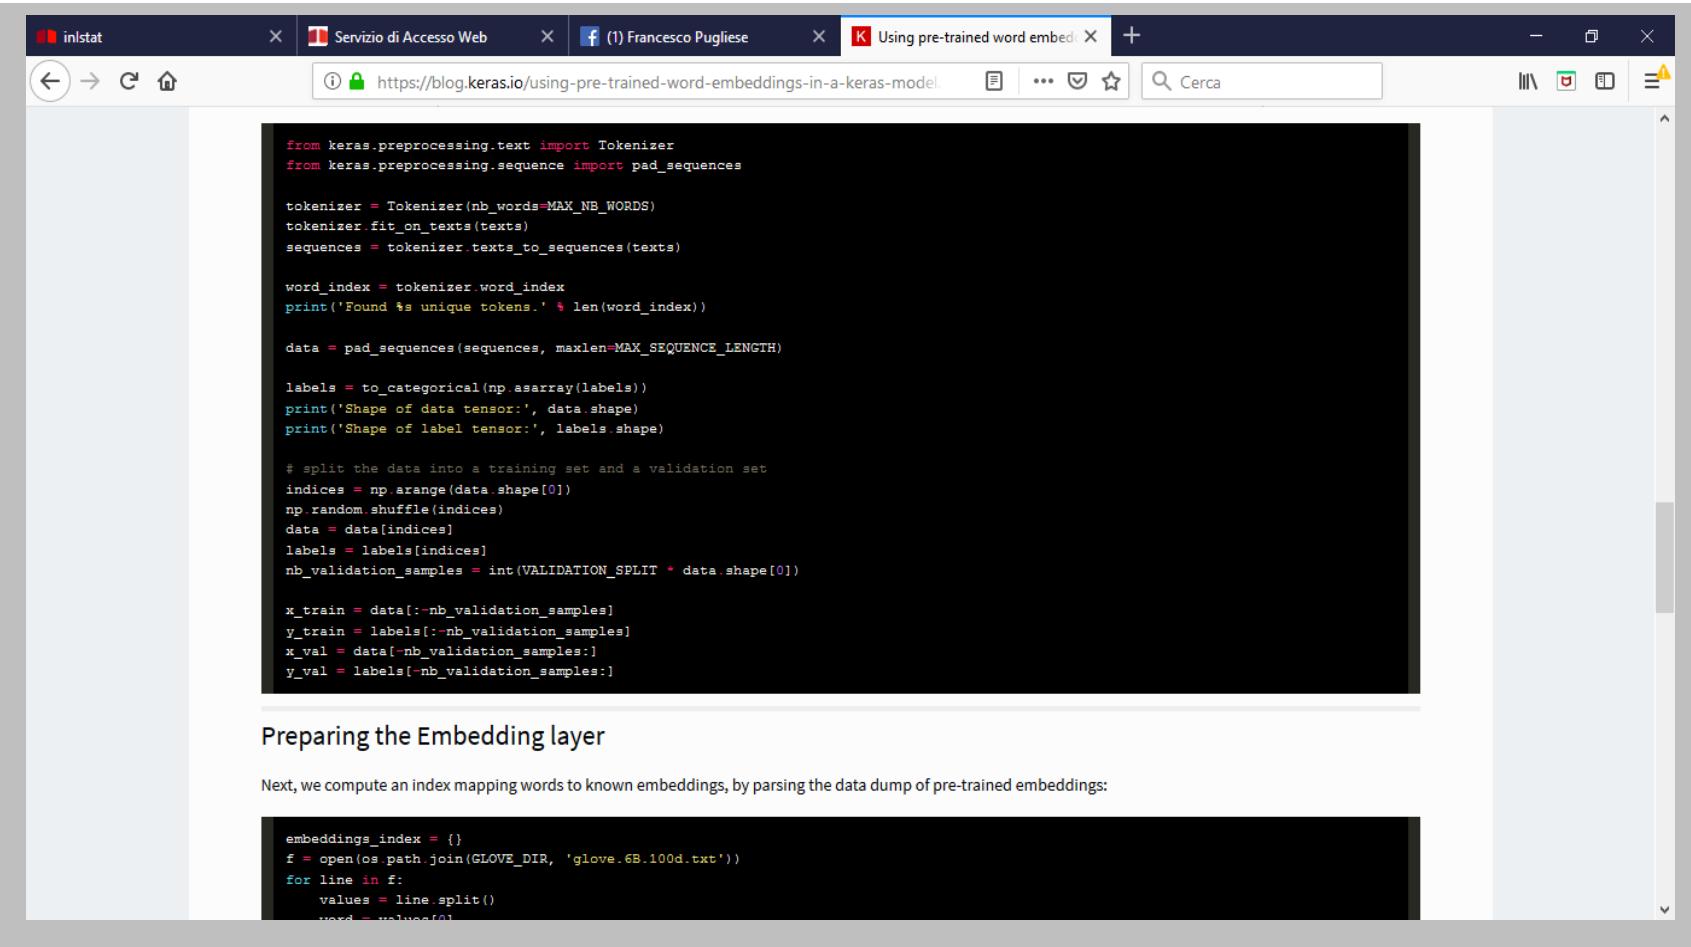
```
texts = [] # list of text samples
labels_index = {} # dictionary mapping label name to numeric id
labels = [] # list of label ids
for name in sorted(os.listdir(TEXT_DATA_DIR)):
    path = os.path.join(TEXT_DATA_DIR, name)
    if os.path.isdir(path):
        label_id = len(labels_index)
        labels_index[name] = label_id
        for fname in sorted(os.listdir(path)):
            if fname.isdigit():
                fpath = os.path.join(path, fname)
                if sys.version_info < (3,):
                    f = open(fpath)
                else:
                    f = open(fpath, encoding='latin-1')
                t = f.read()
                i = t.find('\n\n') # skip header
                if 0 < i:
                    t = t[i:]
                texts.append(t)
                f.close()
                labels.append(label_id)

print('Found %s texts.' % len(texts))
```

Then we can format our text samples and labels into tensors that can be fed into a neural network. To do this, we will rely on Keras utilities `keras.preprocessing.text.Tokenizer` and `keras.preprocessing.sequence.pad_sequences`.

```
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
```

Exercise 1: Text Classifier in Python/Keras



The screenshot shows a web browser window with the URL <https://blog.keras.io/using-pre-trained-word-embeddings-in-a-keras-model.html>. The page content displays a Python script for preparing data for a text classifier using Keras.

```
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences

tokenizer = Tokenizer(nb_words=MAX_NB_WORDS)
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)

word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))

data = pad_sequences(sequences, maxlen=MAX_SEQUENCE_LENGTH)

labels = to_categorical(np.asarray(labels))
print('Shape of data tensor:', data.shape)
print('Shape of label tensor:', labels.shape)

# split the data into a training set and a validation set
indices = np.arange(data.shape[0])
np.random.shuffle(indices)
data = data[indices]
labels = labels[indices]
nb_validation_samples = int(VALIDATION_SPLIT * data.shape[0])

x_train = data[:-nb_validation_samples]
y_train = labels[:-nb_validation_samples]
x_val = data[-nb_validation_samples:]
y_val = labels[-nb_validation_samples:]
```

Preparing the Embedding layer

Next, we compute an index mapping words to known embeddings, by parsing the data dump of pre-trained embeddings:

```
embeddings_index = {}
f = open(os.path.join(GLOVE_DIR, 'glove.6B.100d.txt'))
for line in f:
    values = line.split()
    word = values[0]
```

Exercise 1: Text Classifier in Python/Keras

The screenshot shows a Microsoft Edge browser window with the URL <https://blog.keras.io/using-pre-trained-word-embeddings-in-a-keras-model.html>. The page title is "Using pre-trained word embed". The content is a blog post titled "Preparing the Embedding layer". It includes a code snippet for reading GloVe embeddings from a file and creating a dictionary, followed by instructions on how to use this dictionary to create an embedding matrix, and finally code for initializing an Embedding layer in Keras.

Preparing the Embedding layer

Next, we compute an index mapping words to known embeddings, by parsing the data dump of pre-trained embeddings:

```
embeddings_index = {}
f = open(os.path.join(GLOVE_DIR, 'glove.6B.100d.txt'))
for line in f:
    values = line.split()
    word = values[0]
    coefs = np.asarray(values[1:], dtype='float32')
    embeddings_index[word] = coefs
f.close()

print('Found %s word vectors.' % len(embeddings_index))
```

At this point we can leverage our `embedding_index` dictionary and our `word_index` to compute our embedding matrix:

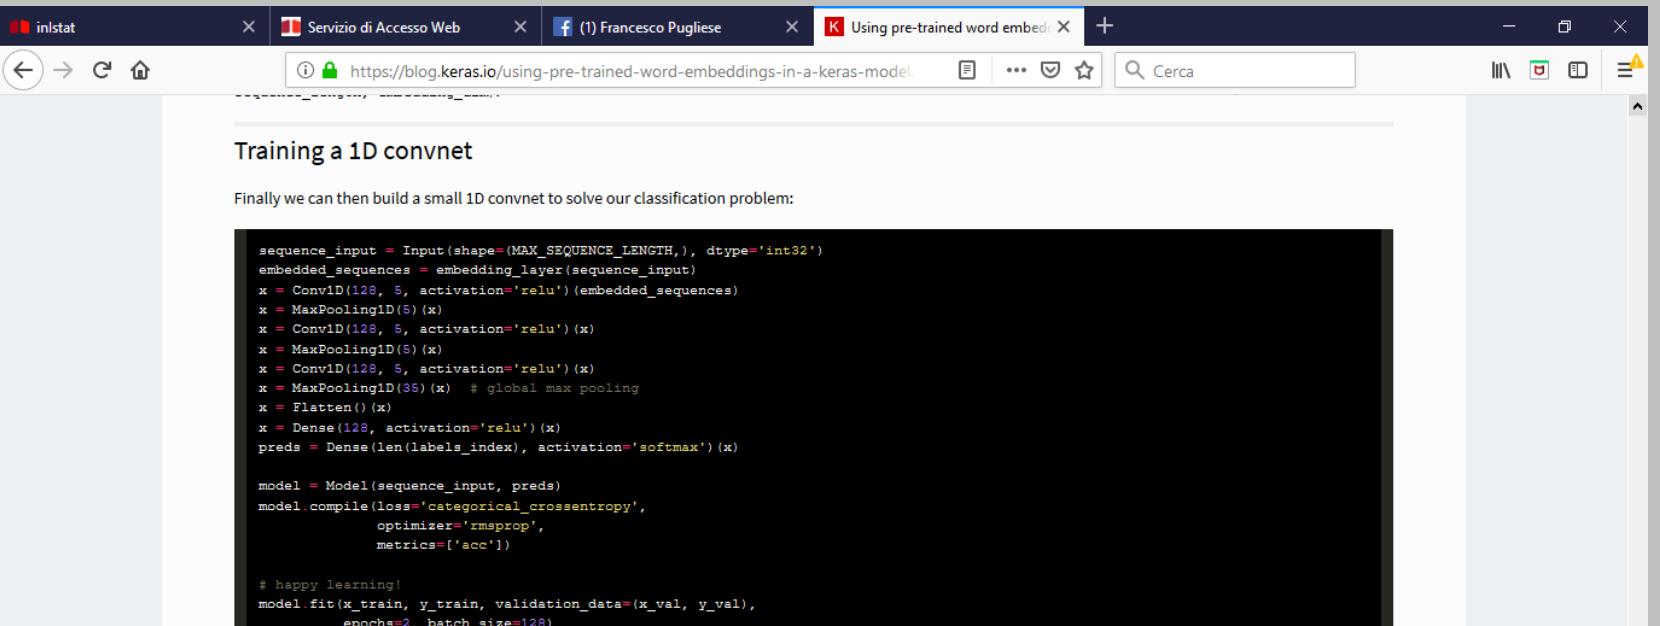
```
embedding_matrix = np.zeros((len(word_index) + 1, EMBEDDING_DIM))
for word, i in word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        # words not found in embedding index will be all-zeros.
        embedding_matrix[i] = embedding_vector
```

We load this embedding matrix into an `Embedding` layer. Note that we set `trainable=False` to prevent the weights from being updated during training.

```
from keras.layers import Embedding

embedding_layer = Embedding(len(word_index) + 1,
                           EMBEDDING_DIM,
                           weights=[embedding_matrix],
                           input_length=MAX_SEQUENCE_LENGTH,
                           trainable=False)
```

Exercise 1: Text Classifier in Python/Keras



The screenshot shows a web browser window with four tabs open:

- instat
- Servizio di Accesso Web
- (1) Francesco Pugliese
- K Using pre-trained word embed... (active tab)

The active tab displays a Python code snippet for training a 1D convolutional neural network (convnet) on text data using pre-trained word embeddings. The code defines an input layer, several convolutional layers with max pooling, and a final dense layer for classification.

```
sequence_input = Input(shape=(MAX_SEQUENCE_LENGTH,), dtype='int32')
embedded_sequences = embedding_layer(sequence_input)
x = Conv1D(128, 5, activation='relu')(embedded_sequences)
x = MaxPooling1D(5)(x)
x = Conv1D(128, 5, activation='relu')(x)
x = MaxPooling1D(5)(x)
x = Conv1D(128, 5, activation='relu')(x)
x = MaxPooling1D(35)(x) # global max pooling
x = Flatten()(x)
x = Dense(128, activation='relu')(x)
preds = Dense(len(labels_index), activation='softmax')(x)

model = Model(sequence_input, preds)
model.compile(loss='categorical_crossentropy',
              optimizer='rmsprop',
              metrics=['acc'])

# happy learning!
model.fit(x_train, y_train, validation_data=(x_val, y_val),
          epochs=2, batch_size=128)
```

This model reaches **95% classification accuracy** on the validation set after only 2 epochs. You could probably get to an even higher accuracy by training longer with some regularization mechanism (such as dropout) or by fine-tuning the Embedding layer.

We can also test how well we would have performed by not using pre-trained word embeddings, but instead initializing our Embedding layer from scratch and learning its weights during training. We just need to replace our Embedding layer with the following:

```
embedding_layer = Embedding(len(word_index) + 1,
                            EMBEDDING_DIM,
                            input_length=MAX_SEQUENCE_LENGTH)
```

REFERENCES

- Scott, A. J., & Knott, M. (1974).** A cluster analysis method for grouping means in the analysis of variance. *Biometrics*, 507-512.
- Bao, W., Yue, J., & Rao, Y.** (2017). A deep learning framework for financial time series using stacked autoencoders and long-short term memory. *PloS one*, 12(7), e0180944.
- Bodyanskiy Y, Popov S.** (2006) Neural network approach to forecasting of quasiperiodic financial time series. *Eur J Oper Res*. 175(3):1357–66.
- Nourani V, Komasi M, Mano A.** (2009) A multivariate ANN-wavelet approach for rainfall-runoff modeling. *Water Resources Management*. 2009;23(14):2877.
- Kim TY, Oh KJ, Kim C, Do JD.** (2004) Artificial neural networks for non-stationary time series. *Neurocomputing*.
- Vinyals, O., & Le, Q. (2015).** A neural conversational model. *arXiv preprint arXiv:1506.05869*.
- Bahdanau, D., Cho, K., & Bengio, Y. (2014).** Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473*.

REFERENCES

- Hsieh TJ, Hsiao HF, Yeh WC.** (2011) Forecasting stock markets using wavelet transforms and recurrent neural networks: An integrated system based on artificial bee colony algorithm. *Applied Soft Computing*. 2011;11(2):2510–25. 61(C):439–47.
- Sutskever I, Hinton GE (2008)**. Deep, narrow sigmoid belief networks are universal approximators. *Neural Computation*. 20(11):2629–36. pmid:18533819
- Roux NL, Bengio Y. (2010)** Deep Belief Networks Are Compact Universal Approximators. *Neural Computation*.22(8):2192–207.
- Bengio Y, Lamblin P, Popovici D, Larochelle H. (2007)** Greedy layer-wise training of deep networks. *Advances in neural information processing systems*. 19:153.
- Ramsey JB. (1999)** The contribution of wavelets to the analysis of economic and financial data. *Philosophical Transactions of the Royal Society B Biological Sciences*. 357(357):2593–606.
- Palangi H, Ward R, Deng L. (2016)** Distributed Compressive Sensing: A Deep Learning Approach. *IEEE Transactions on Signal Processing*. 64(17):4504–18.
- Amari, S. I., Cichocki, A., & Yang, H. H. (1995, December)**. Recurrent neural networks for blind separation of sources. In *Proc. Int. Symp. NOLTA* (pp. 37-42).

REFERENCES

- Chung, J., Gulcehre, C., Cho, K., & Bengio, Y. (2014).** Empirical evaluation of gated recurrent neural networks on sequence modeling. *arXiv preprint arXiv:1412.3555*.
- Hochreiter S, Schmidhuber J.** Long Short-Term Memory (1997). *Neural Computation*. 1997;9(8):1735–80.
pmid:9377276
- .
- Bliemel F.** Theil's (1973) Forecast Accuracy Coefficient: A Clarification. *Journal of Marketing Research*. 10(4):444.
- Kumar, A., Irsoy, O., Ondruska, P., Iyyer, M., Bradbury, J., Gulrajani, I., ... & Socher, R.** (2016, June). Ask me anything: Dynamic memory networks for natural language processing. In *International Conference on Machine Learning* (pp. 1378-1387).
- Bow, C., Hughes, B., & Bird, S. (2003, July).** Towards a general model of interlinear text. In *Proceedings of EMELD workshop* (pp. 11-13).
- Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., & Dean, J. (2013).** Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems* (pp. 3111-3119).

REFERENCES

Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., ... & Bengio, Y. (2014).
Generative adversarial nets. In *Advances in neural information processing systems* (pp. 2672-2680).

Salimans, T., Goodfellow, I., Zaremba, W., Cheung, V., Radford, A., & Chen, X. (2016). Improved techniques for training gans. In *Advances in Neural Information Processing Systems* (pp. 2234-2242).

.

AKNOWLEDGEMENTS

**THANK YOU
FOR YOUR ATTENTION**

Francesco Pugliese